

# Image classification with convolutional neural networks

---

Deep Learning (Methods) Project 1  
Mikołaj Rzepiński & Krzysztof Wolny

# Presentation Plan

---

Overview

---

Learning Rate

---

Batch Size

---

Dropout

---

L1 regularization

---

Augmentation

---

Few-shot learning

---

Ensembling

---

Conclusions

---

# Used Models

We used three different CNNs:

- **ResNet-50** – older, but still strong
- **EfficientNetV2** – faster and smaller, but still very good
- **ConvNeXt** – one of newest models which takes ideas from transformer models

# Methodology Overview

- Dataset: CINIC-10 (10 classes, color images)
- TensorFlow library
- Google Colab (GPU)
- 10 epochs per experiment
- Pretrained weights from ImageNet

# Default Configuration

- Learning Rate = 0.0005,
- Batch Size = 128,
- Dropout = 0.5,
- L1 Regularization = 0.0,
- Data Augmentation = none,
- No Few-shot Learning.

# Hyperparameter Tuning

- Learning Rate - [0.0001, 0.0003, 0.0005, 0.001]
- Batch Size - [32, 64, 128]
- Dropout - [0.3, 0.5, 0.7]
- L1 Regularization - [0.0, 0.0001, 0.0005]

Only one hyperparameter changed at a time



Learning Rate


# Learning Rate

Higher learning rates resulted in better training accuracy, but validation accuracy became unstable for 0.001 rate.


This aligns with the common observation that too large a learning rate leads to overshooting and poor generalization.



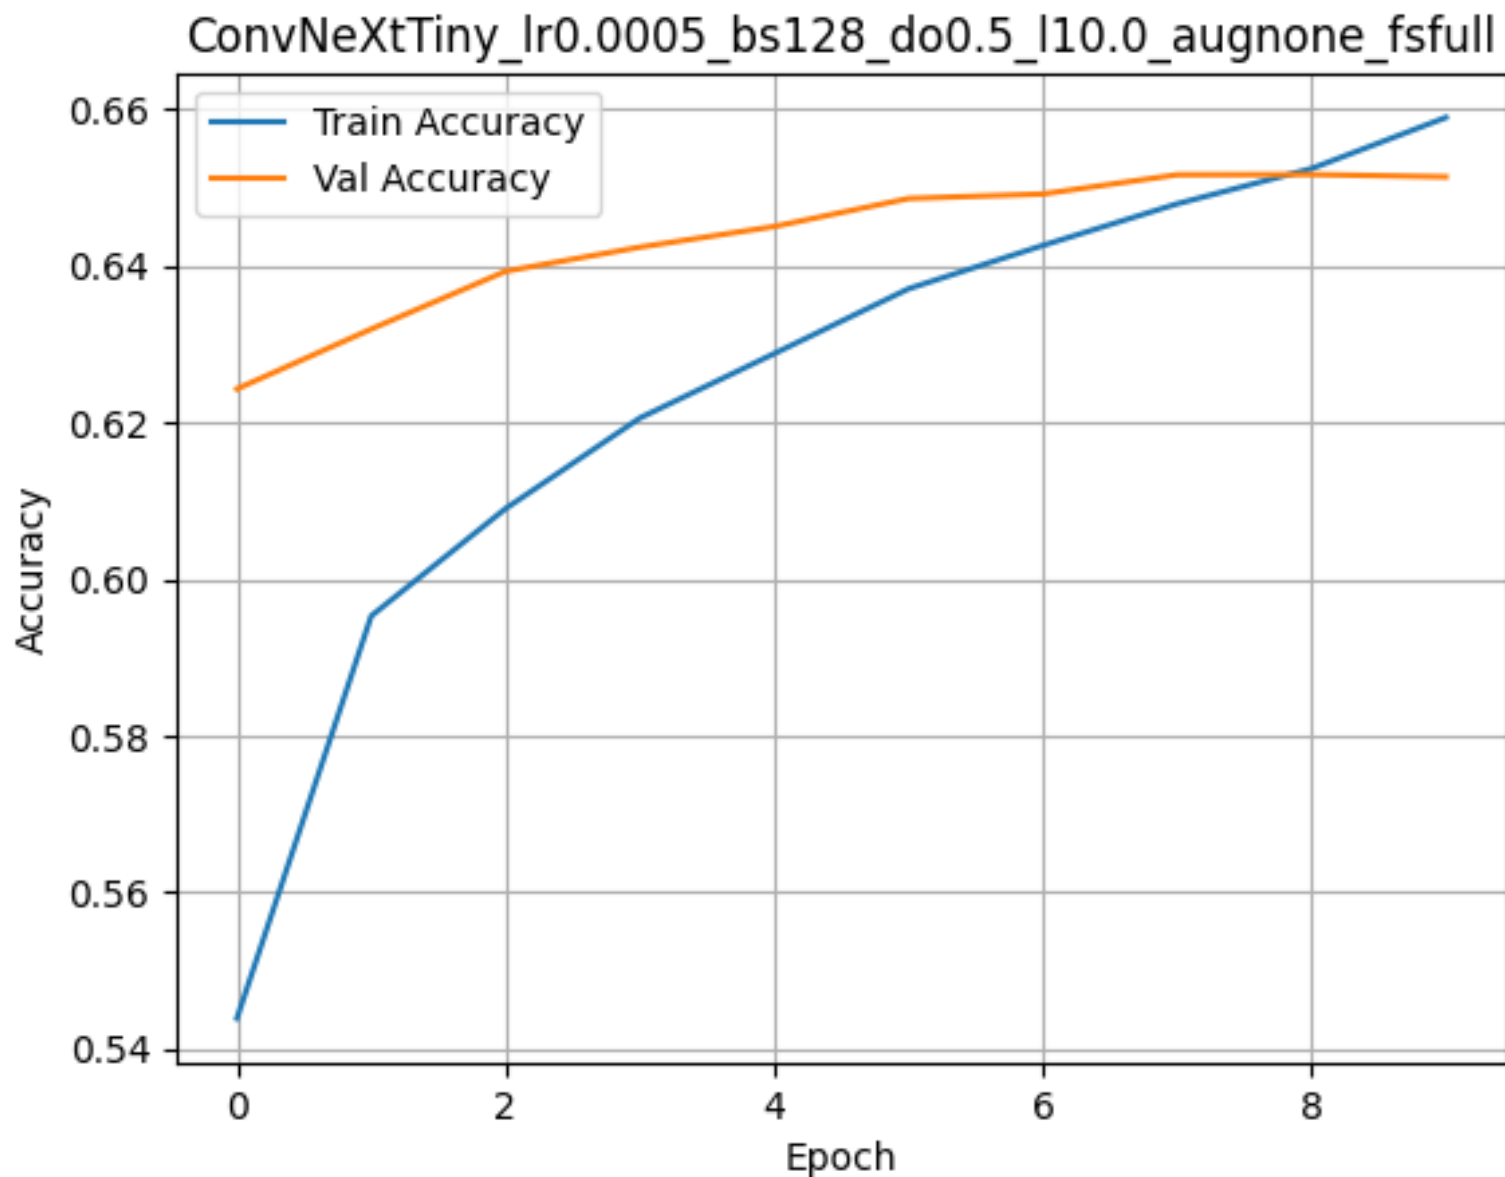
# Learning Rate = 0.0001




# Learning Rate = 0.0003




# Learning Rate = 0.0005



# Learning Rate = 0.001



# Comparison chart for Learning Rate






Batch Size

# Batch Size

As expected, the training and validation performance was relatively insensitive to batch size.


In our setup, changing batch size mostly affected memory usage and iteration speed rather than final accuracy.

# Batch Size = 32







# Batch Size = 64



# Batch Size = 128



# Comparison chart for Batch Size





2,47

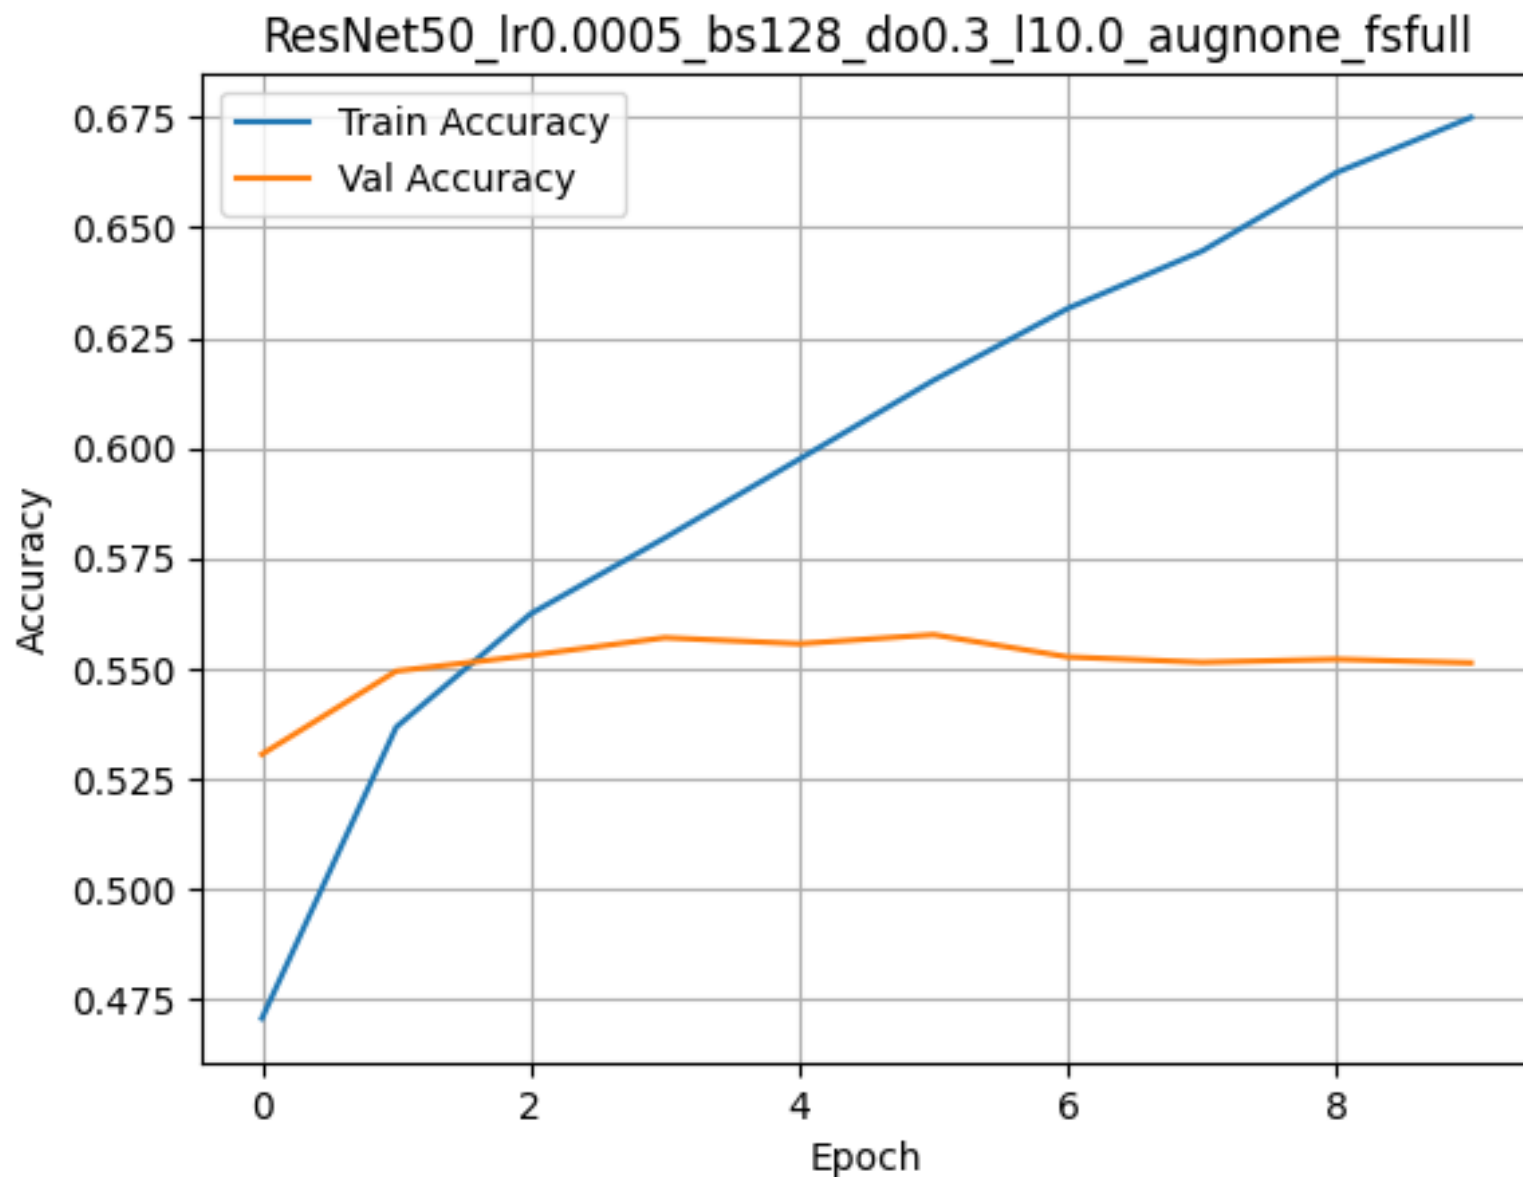
# Dropout

# Dropout


A high dropout rate (0.7) caused clear underfitting, while 0.5 slightly underperformed compared to 0.3.

Given the short training time, aggressive dropout was unnecessary and even made results worse.


# Dropout = 0.3



# Dropout = 0.5




# Dropout = 0.7





# Comparison chart for Dropout





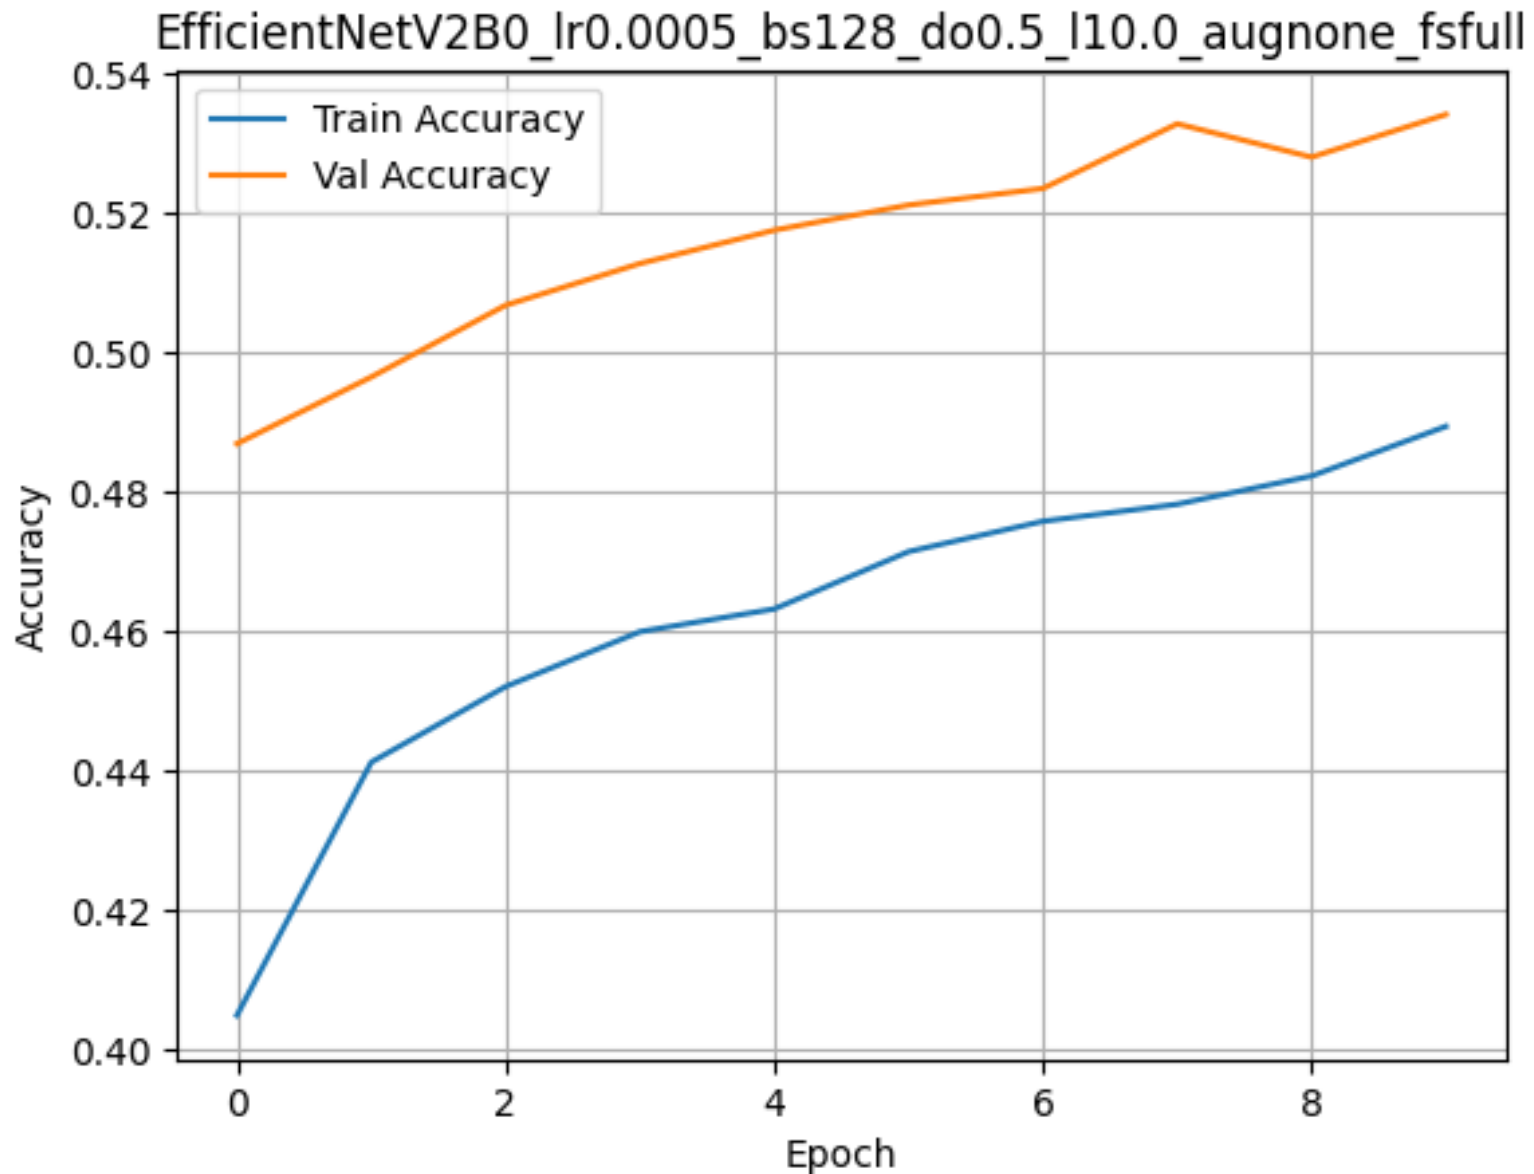
# L1 Regularization

# L1 Regularization


Results indicate that increasing L1 penalty led to decreased accuracy, most notably with 0.0005.

This means that with short training (10 epochs), L1 regularization can limit the model too much and stop it from learning the training data well.


# L1 Regularization = 0.0




# L1 Regularization = 0.0001



# L1 Regularization = 0.0005



# Comparison chart for L1 Regularization






# Augmentation



# Augmentation



- Rotation
- Contrast
- Zoom
- Cutmix



# Augmentation



## Rotation | Contrast

---



# Augmentation

## Zoom | Cutmix






Few-shot  
learning

# Few-shot Learning

- Increasing the number of training images per class (from 10 to 50 to 200) improved model accuracy, with ConvNeXt showing significant gains, especially between 10 and 50 images.
- Using ImageNet weights performed well even with very limited data (10 images per class), demonstrating the power of transfer learning in few-shot learning scenarios.






# Ensembling

# Ensembling

- Hard-voting
- Soft-voting
- Stacking – Logistic Regression







# Conclusions



# Key Results & Observations

- ConvNext got best results out of 3 models we used
- Best Learning Rate: 0.003–0.005
- Batch Size had minor effect (best at 128)
- Lower Dropout (0.3) outperformed higher values
- L1 Regularization hurt short training performance
- CutMix underperformed in short training setup
- Few-shot learning worked even with 10 images/class
- Pretrained models helped generalization
- Best test accuracy: Stacking ensembling 67.85%

# References

- A. F. Agarap, “Deep learning using rectified linear units (relu)”, CoRR, vol. abs/1803.08375, 2018. arXiv: 1803 . 08375. [Online]. Available: [http : // arxiv . org / abs / 1803 .08375](http://arxiv.org/abs/1803.08375).
- K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition”, CoRR, vol. abs/1512.03385, 2015. arXiv:1512 . 03385. [Online]. Available: [http ://arxiv.org/abs/1512.03385](http://arxiv.org/abs/1512.03385).
- M. Tan and Q. V. Le, “Efficientnetv2: Smaller models and faster training”, CoRR, vol. abs/2104.00298, 2021. arXiv:2104.00298. [Online]. Available: <https://arxiv.org/abs/2104.00298>.
- N. Srivastava et al., “Dropout: A simple way to prevent neural networks from overfitting”, Journal of Machine Learning Research, 2014, <https://jmlr.org/papers/v15/srivastava14a.html>.
- C. Shorten and T. M. Khoshgoftaar, “A survey on image data augmentation for deep learning”, Journal of Big Data, vol. 6, no.1, pp. 1–48, 2019. [Online]. Available: <https://arxiv.org/abs/1901.11196>.
- S. Yun et al., “Cutmix: Regularization strategy to train strong classifiers with localizable features”, in ICCV, <https://arxiv.org/abs/1905.04899>, 2019.

Thank you!

